

The Cedilleum Language Specification

Syntax, Typing, Reduction, and Elaboration

Christopher Jenkins

December 12, 2018

1 Introduction

The document describes *Cedilleum*, a general-purpose dependently typed programming language with inductive datatypes. Unlike most languages of this description, the underlying theory of Cedilleum is *not* the Calculus of Inductive Constructions (CIC)[PM15]. Instead, Cedilleum is designed so that it may easily be translated to *Cedille Core* – a compact core theory in which induction is derivable for lambda-encoded datatypes – while still providing high-level features like pattern-matching and recursive definitions. That said, the formal specification of Cedilleum as a self-contained language has a lot in common with CIC – see in particular Section 8 of [Inr18], which served as the basic template for much of this document’s formal development.

1.1 Data-type Declarations

Before diving into the details, let us take a bird’s-eye view of the language by showing some simple example data-type definitions and functions over them.

```
-- Non-recursive
data Bool: * =
  | tt: Bool
  | ff: Bool
.
-- Recursive
data Nat: * =
  | zero: Nat
  | succ: Nat → Nat
.
-- Recursive, parameterized, indexed
data Vec (A: *): Nat → * =
  | vnil : Vec zero
  | vcons: ∀ n: Nat. A → Vec n → Vec (succ n)
.
```

Figure 1: Definition of natural numbers and length-indexed lists

Figure 1 shows some definitions of inductive datatypes, and modulo differences in syntax should seem straightforward to programmers used tilanguages like Agda, Idris, or Coq. Some key differences to note are:

- In constructor type signatures, recursive occurrences of the inductive data-type being defined (such as in `suc : Nat → Nat`) must be positive, *but not strictly positive*.

- In parameterized types (like `Vec` with parameter $(A: \star)$) occurrences of the inductive type being defined are not written applied to its parameters.

For example, the constructor declaration `vnil : Vec zero` results in the term `vnil` having type $\forall A: \star. \text{Vec } A \text{ zero}$ (with \cdot denoting type application)

- In the constructor declaration `vcons : $\forall n: \text{Nat}. A \rightarrow \text{Vec } n \rightarrow \text{Vec } (\text{succ } n)$` , the argument `n` is *computationally irrelevant* (also called *erased*). This is because it is introduced by the irrelevant dependent function former \forall , as opposed to the relevant function former Π . More will be said of this when we discuss the type system of Cedilleum, but for now it suffices to say that implicit quantification comes from the *Implicit Calculus of Constructions*[Miq01].

1.2 Function Definitions

```
-- Non-recursive
ite :  $\forall X: \star. \text{Bool} \rightarrow X \rightarrow X \rightarrow X$ 
  =  $\Lambda X. \lambda b. \lambda \text{then}. \lambda \text{else}. \mu' b \{$ 
    | tt  $\rightarrow \text{then}$ 
    | ff  $\rightarrow \text{else}$ 
   $\}$ .

-- Recursive
add :  $\text{Nat} \rightarrow \text{Nat} \rightarrow \text{Nat}$ 
  =  $\lambda n. \lambda m.$ 
     $\mu \text{rec}. n @(\lambda x: \text{Nat}. \text{Nat}) \{$ 
      | zero  $\rightarrow m$ 
      | succ p  $\rightarrow \text{succ } (\text{rec } p)$ 
     $\}$ .
```

Figure 2: Functions over inductive datatypes

Figure 2 shows functions defined over inductive datatypes using pattern matching and recursion. The first difference to note between the definitions is that functions `ite` and `pred` perform “mere” pattern matching on their arguments by using μ' , whereas `add` uses μ which provides combined pattern-matching and fix-point recursion. In `add`, μ binds `rec` as the name of the fixpoint function for recursion on `n`. From this alone the reader might expect that μ' is merely syntactic sugar for the more verbose μ but without recursion. Actually the difference is a bit more subtle than this, as we will see below in Section 1.3

The first major departure of Cedilleum from other languages with inductive datatypes can be seen in the type of `rec`. The type that the reader might expect it to have is $\Pi x: \text{Nat}. \text{Nat}$ (corresponding to the motive $(\lambda x: \text{Nat}. \text{Nat})$), but in Cedilleum, its type is `rec/type` $\rightarrow \text{Nat}$ (where we read `rec/type` as a single identifier) and by extension for the expression `rec p` to be well-typed, the variable `p` bound in the pattern `succ p` must have type `rec/type`. The name `rec/type` lexically scoped to the body of the μ -expression and is *automatically generated* by Cedilleum by using the name of the recursive function `rec` (which is given by the user) bound by μ . Why introduce this new type? For recursive functions in Cedilleum, *termination is guaranteed by the type system* and not by a separate syntactic check that recursive calls are made on structurally smaller arguments. The type `rec/type` indicates the types of those terms which `rec` may legally take as arguments. These “recursive-occurrence” types appear in the types of constructor sub-data (such as `p` in the example) in patterns introduced by μ , replacing all occurrences of the inductive type itself.

Figure 3 shows the classic dependent function `vappend` over `Vec`, the type of length-indexed lists. Like `add`, it is defined by fixpoint recursion, here over the argument `xs`. Here the fixpoint function `rec` has type $\forall i: \text{Nat}. \Pi zs: \text{rec/type } i. \text{Vec } A (\text{add } i \text{ } n)$, where the recursive occurrence type has kind $\text{Nat} \rightarrow \star$. Note again the missing parameter `A` in the type `rec/type i` – this is not a typo, but rather

```

-- Recursive, parameterized, indexed
vappend : ∀ A: *. ∀ m: Nat. ∀ n: Nat. Vec ·A m → Vec ·A n → Vec ·A (add m n)
= Λ A. Λ m. Λ n. λ xs. λ ys.
  μ rec. xs @ (λ i: Nat. λ zs: Vec ·A i. Vec ·A (add i n)) {
    | vnil          → ys
    | vcons -m' x xs' → [ zs = rec -m' xs' ] - vcons -(add m' n) x zs
  }.

```

Figure 3: Dependent functions over inductive datatypes

an indication that A is “baked-in” to the type `rec/type`. Aside from this the two cases of `vappend` are mostly straightforward: in the `vnil` branch the expected type is `Vec ·A (add zero n)` which converts to `Vec ·A n`, so `ys` suffices; in the `vcons` branch we bind subdata $m': \text{Nat}$, $x: A$, and $xs': \text{rec/type } m'$, with `-m'` indicating that m' is bound *irrelevantly*, then we make a local binding `zs` by invoking recursive function `rec` on m' and xs' (where here `-m'` indicates m' is an irrelevant *argument* to `rec`) before producing a result whose type is convertible with the expected `Vec ·A (add (suc m') n)`.

1.3 Course-of-Value Recursion

We now study Cedilleum’s recursive-occurrence types more closely. Languages with inductive datatypes and recursive function definitions that also wish to have their type systems interpreted as sound logics must address the issue of *termination*, because the principle of general recursion $\forall A: *. (A \rightarrow A) \rightarrow A$ allows one to inhabit every type, the definition of unsoundness! To that end, most such languages perform some termination check separate from type checking that usually involves making sure that arguments to recursive calls are getting *structurally smaller* to ensure that eventually a base case is reached. This check is necessarily conservative (i.e. it will not accept all terminating functions), and the classic example of a function that is not “obviously” terminating is division on natural numbers by iterated subtraction. Intuitively, we understand that subtracting n from m never produces a number larger than m – but it can be tricky to explain this to the termination checker! One advantage that Cedilleum’s type-guided termination checking has is that it allows for a very natural definition of division as iterated subtraction that is “obviously” terminating.

The definition of division is given in Figure 4. Our first definition, `pred'`, is crucial for defining `divide` further below. The type `Rec ·Nat ·R` in its type signature is the type of *witnesses* that terms of type R can be pattern-matched upon as if they had type `Nat`; in the definition of `pred'`, this witness is given name `recR`. All such witnesses are introduced in only one of two ways: once *globally* for each defined data-type (like `Nat` itself – in the definition of `pred` term `Nat/mu` has type `Rec ·Nat ·Nat`; and *locally* for each recursive-occurrence type introduced by μ (in the body of `divide`, term `rec/mu` has type `Rec ·Nat ·rec/type`). In the definition of `pred'`, the notation $\mu' \langle \text{recR} \rangle$ indicates that the witness `recR` is given explicitly to enable (mere) pattern-matching on argument r . After this, the definition of `pred` is easy – it is an instance of `pred'` where R is specialized to `Nat` itself with `Nat/mu`.

Next we define `minus'`. One intuition for understanding the type signature of `minus'` (and `pred'` before it) is that it says “this function will never increase the size of its R argument”. That is to say, to return a term of type R , `minus'` can only return its argument or some sub-data produced by destructing it through pattern-matching. In `pred'` this is done only once; in `minus'` it is done n times by recursion on subtrahend n , invoking `pred'` each time.

Finally, we turn to the definition of `divide` itself. At a high level, we recurse on dividend n , and in the step case test whether n (now known to have the form `succ n'`) is less than the divisor d : if it is we return `zero`, otherwise we recurse after subtracting. First, consider that `succ` has type `Nat → Nat`, but here it is applied to argument n' of type `rec/type`. The reason that this is well-typed is that recursive-occurrence types are always *subtypes* of the inductive data-type from which they came – here, the type of n' is promoted to `Nat` in the application `succ n'`. Second, as mentioned before recursive function `rec` requires an argument of type `rec/type`, so must use `minus'` (passing the generated, locally bound witness

```

pred' : ∀ R: *. Rec ·Nat ·R ⇒ R → R
  = Λ R. Λ recR. λ r. μ'<recR> r { | zero → r | succ r' → r' }.

pred : Nat → Nat = pred' -Nat/mu .

minus' ∀ R: *. Rec ·Nat ·R ⇒ R → Nat → R
  = Λ R. Λ recR. λ m. λ n. μ rec. n @(λ _: Nat. R) {
    | zero    → m
    | succ n' → pred' -recR (rec n')
  }.

minus : Nat → Nat → Nat = minus' -Nat/Rec .

lt : Nat → Nat → Bool
  = λ m. λ n. μ' (minus m n) {
    | zero    → ff
    | succ _ → tt
  }.

divide : Nat → Nat → Nat
  = λ n. λ d. μ rec. n @(λ _: Nat. Nat) {
    | zero    → zero
    | succ n' →
      ite (lt (succ n') d) zero
        (succ (rec (minus' -rec/mu n' (pred d))))
  }.

```

Figure 4: Histomorphic recursion and division

`rec/mu`) so that after subtraction we have a value of the appropriate type. Finally, note how this definition of division is terminating even when divisor `d` is `zero` – in that case the expression `minus' -rec/mu n' (pred d)` reduces to `n'`, which is *still* smaller than the expression this recursive case started with, `succ n'`. Put another way: the “naive” expression `rec (minus (succ n') d)` which is potentially non-terminating is not accepted by Cedilleum *because it is ill-typed*.

1.4 Reasoning via Induction

```

add-zero-r : Π m: Nat. {add m zero ≃ m}
  = λ m. μ ih. m @(λ x: Nat. {add x zero ≃ x}) {
    | zero    → β
    | succ r → χ {succ (add r zero) ≃ succ r} - ρ (ih r) - β
  } .

```

Figure 5: A proof via induction

Figure 5 shows a simple proof that `zero` is the right identity of `add` using induction on `Nat`. In the base case, pattern `zero` is substituted in for `x` in the motive, and the expected result type of the branch is `{add zero zero ≃ zero}`, which is true by reflexivity (notated `β`) after conversion. In the step case, pattern `succ r` is substituted in for `x` in the motive (which is well-kinded in general thanks to subtyping, as `succ r` has type `Nat`, even though `r` has type `ih/type`), and the expected result type of the branch is `{add (succ r`

$\text{zero} \simeq \text{succ } r$. Operator χ allows users to write type annotations, so the expected type converts to the annotated type $\{\text{succ } (\text{add } r \text{ zero}) \simeq \text{succ } r\}$, which is then rewritten using the inductive hypothesis $\text{ih } r$.

1.5 Reduction Rules of μ and μ'

In Section 1.4 there was an omission regarding convertibility of terms in the expected type of case branches. For example, the expected type corresponding to the branch $\text{succ } r$ in the definition of **add-zero-r** above is $\{\text{add } (\text{succ } r) \text{ zero} \simeq \text{succ } r\}$. By β -reduction alone, this reduces to

```
{  $\mu$  rec. (succ r) {
  | zero  $\rightarrow$  zero
  | succ p  $\rightarrow$  succ (rec p)
}  $\simeq$  succ r
}
```

To get the left-hand side of this equation to be convertible with $\text{succ } (\text{add } r \text{ zero})$, we need μ -reduction rules. μ -reduction is a combination of fix-point unrolling and is typically called δ -reduction for languages with inductive data-types. Here, because the scrutinee is $\text{succ } r$, then entire μ -expression reduces to the body of the case-branch guarded by $\text{succ } p$, with recursive function **rec** replaced by the entire μ -expression itself. Thus, the equation above reduces to

```
{ succ ( $\mu$  rec. r {
  | zero  $\rightarrow$  zero
  | succ p  $\rightarrow$  succ (rec p)
})  $\simeq$  succ r
}
```

where the left-hand side is now convertible with $\text{succ } (\text{add } r \text{ zero})$. μ' -reduction works similarly, except there is no fixpoint to unroll.

1.6 Non-strictly Positive Datatypes

In the preceeding sections, we have that seen “cast” functions like Nat/ih (in Figure 5) show up in the the expected type of a case branch, and also have noted already that Cedilleum allows for positive but not strictly positive data type definitions. We now take a look at how these two things interact.

Figure 6 presents a definition of **PTree**, an infinitary tree which is not strictly positive in the **node** constructor, and a proof of induction for it using μ . One intuition for what kind of terms inhabit **PTree** is “at a **node**, there must be some way of selecting one **PTree** (of infinitely many) from some partition $\text{PTree} \rightarrow \text{Bool}$ ”. The branch given by pattern **leaf** is the **base** case, requiring a proof of $P \text{ leaf}$ which we have by assumption. For the **step** case given by the branch for pattern **node**, the expected type is the trickier $P (\text{node } s')$, where s' is defined as above. Recall that in our branch the constructor arguments have all recursive occurrences of their inductive type replaced with a special “abstracted” version. In the case of **leaf**, the subdata s has type $(\text{PTree/ih} \rightarrow \text{Bool}) \rightarrow \text{PTree/ih}$. The the naive expected branch type $P (\text{node } s)$ given by simply substituting the pattern in for the bound x given in the motive is not well-kinded! To fix this we η -expand and cast using **fromPTree/ih** as needed to produce the appropriate expected type.

Now we examine the body of the **node** branch itself more closely. First, we need a way to convert any “partition” p of type $\text{PTree} \rightarrow \text{Bool}$ to $\text{PTree/ih} \rightarrow \text{Bool}$ so that it can operate over terms of the abstracted type. With this we can do the same for our selector s , producing s' . Finally, we invoke **step** so that we can have prove $P (\text{node } s')$ as desired. The second argument of **step** requires a proof that $P (s' p)$ holds for any p . To show this, we take the assumed p , “weaken” it to work over the abstracted recursive types with **conv-p**, and pass this to s (and not s' !) to produce a value of type PTree/ih suitable for consumption by **ih**, which has type $\Pi t: \text{PTree/ih}. P (\text{fromPTree/ih } t)$. Note again that the proof

```

data PTree : * =
  | leaf : PTree
  | node : ((PTree → Bool) → PTree) → PTree
.
PTreeSel : * = (PTree → Bool) → PTree.
indTree : ∀ P: PTree → *.
  P leaf → (∀ s: PTreeSel. (Π p: PTree → Bool. P (s p)) → P (node s))
  → Π t: PTree. P t
= Λ P. λ base. λ step. λ t. μ ih. t @ (λ x: PTree. P x) {
  | leaf → base
  | node s →
    [ conv-p : (PTree → Bool) → PTree/ih → Bool
      = λ p. λ r. p (fromPTree/ih r) ]
    - [ s' : PTreeSel = λ p. fromPTree/ih (s (conv-p p)) ]
    - step -s' (λ p. ih (s (conv-p p)))
  }.

```

Figure 6: A non-strictly positive infinitary tree

we need to give here, $P (s' p)$, is convertible with the proof $P (fromPTree/ih (s (conv-p p)))$ that the call to ih actually returns.

2 Syntax

id	identifiers for definitions
u, c	term variables
X	type variables
κ	kind variables
$x ::= id \mid u \mid X$	non-kind variables
$y ::= x \mid \kappa$	all variables

Figure 7: Identifiers

Identifiers We now turn to a more formal treatment of Cedilleum. Figure 7 gives the metavariables used in our grammar for identifiers. We consider all identifiers as coming from two distinct lexical “pools” – regular identifiers (consisting of identifiers id given for modules and definitions, term variables u , and type variables X) and kind identifiers κ . In Cedilleum source files (as in the parent language Cedille) kind variables should be literally prefixed with κ – the suffix can be any string that would by itself be a legal non-kind identifier. For example, `myDef` is a legal term / type variable and a legal name for a definition, whereas `κmyDeff` is only legal as a kind definition.

Untyped Terms The grammar of pure (untyped) terms the untyped λ -calculus augmented with a primitives for combination fixed-point and pattern-matching definitions (and an auxiliary pattern-matching construct).

Modules and Definitions All Cedilleum source files start with production *mod*, which consists of a module declaration, a sequence of import statements which bring into scope definitions from other source files, and a sequence of *commands* defining terms, types, and kinds. As an illustration, consider the first few lines of a hypothetical `list.ced`:

f, p	$::=$	u, v, c	variables
		$\lambda u. p$	functions
		$f p$	applications
		$\mu u . p \{pcase^*\}$	fixed-point and pattern matching
		$\mu' p \{pcase^*\}$	simple pattern matching
$pcase$	$::=$	$ u u^* \mapsto f$	

Figure 8: Untyped terms

mod	$::=$	module id . $imprt^* cmd^*$	module declarations
$imprt$	$::=$	import id .	module imports
cmd	$::=$	$defTermOrType$ $defDataType$ $defKind$	definitions
$defTermOrType$	$::=$	$id \text{ checkType}^? = t$.	term definition
		$id : K = T$.	type definition
$defKind$	$::=$	$\kappa = K$	kind definition
$defDataType$	$::=$	data $id \text{ param}^* : K = constr^*$.	datatype definitions
$checkType$	$::=$	$: T$	annotation for term definition
$param$	$::=$	$(x : C)$	
$constr$	$::=$	$ id : T$	

Figure 9: Modules and definitions

```
module list .
```

```
import nat .
```

Imports are handled first by consulting a global options files known to the Cedilleum compiler (on *nix systems `~/cedille/options`) containing a search path of directories, and next (if that fails) by searching the directory containing the file being checked.

Term and type definitions are given with an identifier, a classifier (type or kind, resp.) to check the definition against, and the definition. For term definitions, giving classifier (i.e. the type) is optional. As an example, consider the definitions for the type of Church-encoded lists and two variants of the nil constructor, the first with a top-level type annotation and the second with annotations sprinkled on binders:

```
cList : * → *
  = λ A : * . ∀ X : * . (A → X → X) → X → X .

cNil  : ∀ A : * . cList · A
  = λ A . λ X . λ c . λ n . n .
cNil' = λ A : * . λ X : * . λ c : A → X → X . λ n : X . n .
```

Kind definitions are given without classifiers (all kinds have super-kind \square), e.g. $\kappa\text{func} = * \rightarrow *$

Inductive datatype definitions take a set of *parameters* (term and type variables which remain constant throughout the definition) well as a set of *indices* (term and type variables which *can* vary), followed by zero or more constructors. Each constructor begins with “|” (though the grammar can be relaxed so that the

first of these is optional) and then an identifier and type is given. As an example, consider the following two definitions for lists and vectors (length-indexed lists).

```

data Bool :  $\star$  =
  | tt : Bool
  | ff : Bool
.
data Nat :  $\star$  =
  | zero : Nat
  | suc  : Nat  $\rightarrow$  Nat
.
data List (A :  $\star$ ) :  $\star$  =
  | nil  : List
  | cons : A  $\rightarrow$  List  $\rightarrow$  List
.
data Vec (A :  $\star$ ) : Nat  $\rightarrow$   $\star$  =
  | vnil  : Vec zero
  | vcons :  $\forall$  n: Nat. A  $\rightarrow$  Vec n  $\rightarrow$  Vec (succ n)
.

```

Sorts S	$::=$	\square	sole super-kind
		K	kinds
Classifiers C	$::=$	K	kinds
		T	types
Kinds K	$::=$	$\Pi x : C . K$	explicit product
		$C \rightarrow K$	kind arrow
		\star	the kind of types that classify terms
		(K)	
Types T	$::=$	$\Pi x : T . T$	explicit product
		$\forall x : C . T$	implicit product
		$\lambda x : C . T$	type-level function
		$T \Rightarrow T'$	arrow with erased domain
		$T \rightarrow T'$	normal arrow type
		$T \cdot T'$	application to another type
		$T \ t$	application to a term
		$\{ p \simeq p' \}$	untyped equality
		(T)	
		X	type variable
		\bullet	hole for incomplete types

Figure 10: Kinds and types

Types and Kinds In Cedilleum, the expression language is stratified into three main “classes”: kinds, types, and terms. Kinds and types are listed in Figure 10 and terms are listed in Figure 11 along with some auxiliary grammatical categories. In both of these figures, the constructs forming expressions are listed from lowest to highest precedence – “abstractors” ($\lambda \ \Lambda \ \Pi \ \forall$) bind most loosely and parentheses most tightly. Associativity is as-expected, with arrows ($\rightarrow \Rightarrow$) and applications being left-associative and abstractors being right-associative.

The language of kinds and types is similar to that found in the Calculus of Implicit Constructions¹. Kinds are formed by dependent and non-dependent products (Π and \rightarrow) and a base kind for types which can classify terms (\star). Types are also formed by the usual (dependent and non-dependent) products (Π and \rightarrow) and also *implicit* products (\forall and \Rightarrow) which quantify over erased arguments (that is, arguments that disappear at run-time). Π -products are only allowed to quantify over terms as all types occurring in terms are erased at run-time, but \forall -products can quantify over types *and* terms because terms can be erased. Meanwhile, non-dependent products (\rightarrow and \Rightarrow) can only “quantify” over terms because non-dependent type quantification does not seem particularly useful. Besides these, Cedilleum features type-level functions and applications (with term and type arguments), and a primitive equality type for untyped terms. Last of all is the “hole” type (\bullet) for writing partial type signatures or incomplete type applications. There are term-level holes as well, and together the two are intended to help facilitate “hole-driven development”: any hole automatically generates a type error and provides the user with useful contextual information.

We illustrate with another example: what follows is a module stub for **DepCast** defining dependent casts – intuitively, functions from $a : A$ to B that are also equal² to identity – where the definitions **CastE** and **castE** are incomplete.

```
module DepCast .
```

```
CastE <|  $\Pi A : \star . (A \rightarrow \star) \rightarrow \star = \bullet$  .
```

```
castE <|  $\forall A : \star . \forall B : A \rightarrow \star . \text{CastE} \cdot A \cdot B \Rightarrow \Pi a : A . B a = \bullet$  .
```

Annotated Terms Terms can be explicit and implicit functions (resp. indicated by λ and Λ) with optional classifiers for bound variables, let-bindings, applications $t\ t'$, $t\ -t'$, and $t\ \cdot T$ (resp. to another term, an erased term, or a type). In addition to this there are a number of useful operators for equational reasoning, type casting, providing annotations, and pattern matching. Each operator will be discussed in more detail in Section 4, but a few concrete programs in Cedilleum are given below merely to give a better idea of the syntax of the language.

```
isvnil :  $\forall A : \star . \forall n : \text{Nat} . \text{Vec } A\ n \rightarrow \text{Bool}$ 
      =  $\Lambda A . \Lambda n . \lambda xs . \mu' xs\ @(\Lambda n . \lambda xs . \text{Bool}) \{$ 
          | vnil           $\rightarrow \text{tt}$ 
          | vcons -n x xs  $\rightarrow \text{ff}$ 
         $\}$ .
vlength :  $\forall A : \star . \forall n : \text{Nat} . \text{Vec } A\ n \rightarrow \text{Nat}$ 
      =  $\Lambda A . \Lambda n . \lambda xs . \mu\ \text{len} . xs\ @(\Lambda n . \lambda x . \text{Nat}) \{$ 
          | vnil           $\rightarrow \text{zero}$ 
          | vcons -n x xs  $\rightarrow \text{suc } (\text{len } -n\ xs)$ 
         $\}$ .
```

3 Erasure and Reduction

The definition of the erasure function given in Figure 12 takes the annotated terms from Figures 10 and 11 to the untyped terms of Figure 8. The last two equations indicate that any type or erased arguments in the zero or more *vararg*’s of pattern-match case are indeed erased. The additional constructs introduced in the annotated term language such as β , ϕ , and ρ , are all erased to the language of pure terms.

Reduction rules are defined for the untyped term language. In essence, to run a Cedilleum program you first erase it, then reduce it.

¹Cite

²Module erasure, discussed below

Subjects s	$::= t$	term
	T	type
Terms t	$::= \lambda x \text{ class}^?. t$	normal abstraction
	$\Lambda x \text{ class}^?. t$	erased abstraction
	$[\text{defTermOrType}] - t$	let definitions
	$\rho \ t - t'$	equality elimination by rewriting
	$\phi \ t - t' \ \{t''\}$	type cast
	$\chi \ T - t$	check a term against a type
	$\delta - t$	ex falso quodlibet
	$\theta \ t \ t'^*$	elimination with a motive
	$t \ t'$	applications
	$t \ -t'$	application to an erased term
	$t \ .T$	application to a type
	$\beta \ \{t\}$	reflexivity of equality
	$\varsigma \ t$	symmetry of equality
	$\mu \ u \ . \ t \ \text{motive}^? \ \{\text{case}^*\}$	type-guarded pattern match and fixpoint
	$\mu' \ t \ \text{motive}^? \ \{\text{case}^*\}$	auxiliary pattern match
	u	term variable
	(t)	
	\bullet	hole for incomplete term
case	$::= \mid c \ \text{vararg}^* \rightarrow t$	pattern-matching cases
vararg	$::= u$	normal constructor argument
	$-u$	erased constructor argument
	$\cdot X$	type constructor argument
class	$::= : C$	
motive	$::= @ \ T$	motive for induction

Figure 11: Annotated Terms

β -reduction

$$(\lambda x. p_1) \ p_2 \rightsquigarrow_{\beta} [p_2/x]p_1$$

The rule for β -reduction is standard: those expressions consisting of a λ -abstraction as the left component of an application reduce by having their bound variable substituted away by the given argument (where $[p_2/x]$ is the simultaneous and capture-avoiding substitution of p_2 for x)

μ' -reduction

$$\mu' \ (c_i \ p_1 \dots p_n) \ \{\dots \mid c_i \ u_1 \dots u_n \mapsto f \mid \dots\} \rightsquigarrow_{\mu'} [p_1 \dots p_n / u_1 \dots u_n] f$$

μ' -reduction is a simple pattern-matching reduction rule: if the scrutinee of μ' is some variable-headed application $c_i \ p_1 \dots p_n$ where the head c_i matches one of the branch patterns, replace the entire expression with the branch body f after substituting each of the bound variables of the branch pattern $u_1 \dots u_n$ with the scrutinee's arguments $p_1 \dots p_n$

μ -reduction

$$\frac{\exists i. \ c = c_i \wedge j_i = n \quad p_{\mu} = \lambda v. \mu \ u. \ v \ \{c_i \ u_{i1} \dots u_{ij_i} \mapsto f_i\}_{i=1..n}}{\mu \ u. (c \ p_1 \dots p_n) \ \{c_i \ u_{i1} \dots u_{ij_i} \mapsto f_i\}_{i=1..n} \rightsquigarrow_{\mu} [p_1 \dots p_n / u_1 \dots u_n] [u/p_{\mu}] f} \ \mu$$

μ -reduction is similar to μ' -reduction, but combines with it fixpoint reduction. Again, if the scrutinee $c \ p_1 \dots p_n$ matches one of the branch patterns $c_i \ u_{i1} \dots u_{ij_i}$ (for some i , where $j_i = n$), then we replace the original μ expression with the matched branch, replacing each of the pattern variables $u_1 \dots u_n$ with the

$ x $	$= x$
$ \star $	$= \star$
$ \square $	$= \square$
$ \beta \{t\} $	$= t $
$ \delta t $	$= t $
$ \chi T^? - t $	$= t $
$ \varsigma t $	$= t $
$ t t' $	$= t t' $
$ t - t' $	$= t $
$ t \cdot T $	$= t $
$ \rho t - t' $	$= t' $
$ \forall x:C. C' $	$= \forall x: C . C' $
$ \Pi x:C. C' $	$= \Pi x: C . C' $
$ \lambda u:T. t $	$= \lambda u. t $
$ \lambda u. t $	$= \lambda u. t $
$ \lambda X:K. C $	$= \lambda X: K . C $
$ \Lambda x:C. t $	$= t $
$ \phi t - t' \{t''\} $	$= t'' $
$ [x = t : T] - t' $	$= (\lambda x. t') t $
$ [X = T : K] - t $	$= t $
$ \{t \simeq t'\} $	$= \{ t \simeq t' \}$
$ \mu u. \cdot t \text{ motive}^? \{case^*\} $	$= \mu u. t \{ case^* \}$
$ \mu' t \text{ motive}^? \{case^*\} $	$= \mu' t \{ case^* \}$
$ id \text{ vararg}^* \mapsto t $	$= id vararg^* \mapsto t $
$ -u $	$=$
$ \cdot X $	$=$

Figure 12: Erasure for annotated terms

scrutinee's arguments $p_1 \dots p_n$, but *in addition* we also replace the μ -bound variable u (which represents the entire μ expression itself) with a function p_μ that takes its argument v and re-creates the original μ expression by scrutinizing v .

4 Type System (sans Inductive Datatypes)

Figure 13: Contexts

Typing contexts $\Gamma ::= \emptyset \mid x:C, \Gamma \mid x=s:C, \Gamma$

The inference rules for classifying expressions in Cedilleum are stratified into two judgments. Figure 14 gives the uni-directional rules for ensuring types are well-kinded and kinds are well-formed. Future versions of Cedilleum will allow for bidirectional checking for both typing *and* sorting, allowing for a unification of these two figures. Most of these rules are similar to what one would expect from the Calculus of Implicit Constructions, so we focus on the typing rules unique to Cedilleum.

⁴Where we assume t does not occur anywhere in T

⁴Where $\mathbf{tt} = \lambda x. \lambda y. x$ and $\mathbf{ff} = \lambda x. \lambda y. y$

$$\begin{array}{c}
\overline{\Gamma \vdash \star : \square} \\
\\
\frac{FV(p \ p') \subseteq \text{dom}(\Gamma)}{\Gamma \vdash \{p \simeq p'\} : \star} \quad
\frac{\Gamma \vdash C : S \quad \Gamma, y : C \vdash C' : S'}{\Gamma \vdash \Pi y : C. C' : S'} \quad
\frac{\Gamma \vdash C : S \quad \Gamma, y : C \vdash C' : \star}{\Gamma \vdash \forall y : C. C' : \star} \\
\\
\overline{\Gamma \vdash \kappa : \Gamma(\kappa)} \quad
\overline{\Gamma \vdash X : \Gamma(X)} \\
\\
\frac{\Gamma \vdash \Pi x : C. K : \square \quad \Gamma, x : C \vdash T : K}{\Gamma \vdash \lambda x : C. T : \Pi x : C. K} \quad
\frac{\Gamma \vdash T : \Pi x : K. K' \quad \Gamma \vdash T' : K}{\Gamma \vdash T \cdot T' : [T'/x]K'} \quad
\frac{\Gamma \vdash T : \Pi x : T'. K \quad \Gamma \vdash_{\Downarrow} t : T'}{\Gamma \vdash T t : [t/x]K}
\end{array}$$

Figure 14: Sort checking $\boxed{\Gamma \vdash C : S}$

$$\begin{array}{c}
\overline{\Gamma \vdash_{\delta} u : \Gamma(u)} \quad
\frac{\Gamma \vdash T : K \quad \Gamma, x : T \vdash_{\delta} t : T'}{\Gamma \vdash_{\delta} \lambda x : T. t : \Pi x : T. T'} \quad
\frac{\Gamma, x : T \vdash_{\Downarrow} t : T'}{\Gamma \vdash_{\Downarrow} \lambda x. t : \Pi x : T. T'} \\
\\
\frac{\Gamma \vdash C : S \quad x \notin FV(|t|) \quad \Gamma, x : C \vdash_{\delta} t : T}{\Gamma \vdash_{\delta} \Lambda x : C. t : \forall x : C. T} \quad
\frac{x \notin FV(|t|) \quad \Gamma, x : C \vdash_{\delta} t : T}{\Gamma \vdash_{\Downarrow} \Lambda x. t : \forall x : C. T} \quad
\frac{\Gamma \vdash_{\Uparrow} t : \Pi x : T'. T \quad \Gamma \vdash_{\Downarrow} t' : T'}{\Gamma \vdash_{\delta} t t' : [t'/x]T} \\
\\
\frac{\Gamma \vdash_{\Uparrow} t : \forall X : K. T' \quad \Gamma \vdash T : K}{\Gamma \vdash_{\delta} t \cdot T : [T/X]T'} \quad
\frac{\Gamma \vdash_{\Uparrow} t : \forall x : T'. T \quad \Gamma \vdash_{\Downarrow} t' : T'}{\Gamma \vdash_{\delta} t \cdot t' : [t'/x]T} \quad
\frac{\Gamma \vdash_{\Uparrow} t : T' \quad |T'| =_{\beta} |T|}{\Gamma \vdash_{\Downarrow} t : T} \\
\\
\frac{\Gamma \vdash T : K \quad \Gamma \vdash_{\Downarrow} t : T \quad \Gamma, id = t : T \vdash_{\delta} t' : T'}{\Gamma \vdash_{\delta} [id : T = t] \cdot t' : T'} \quad
\frac{\Gamma \vdash_{\Uparrow} t : T \quad \Gamma, id = t : T \vdash_{\delta} t' : T'}{\Gamma \vdash_{\delta} [id = t] \cdot t' : T'} \quad
\frac{\Gamma \vdash_{\Uparrow} t : \{t_1 \simeq t_2\} \quad \Gamma \vdash_{\Uparrow} t' : [t_1/x] T}{\Gamma \vdash_{\delta} \rho t \cdot t' : [t_2/x] T} \quad 3 \\
\\
\frac{\Gamma \vdash K : \square \quad \Gamma \vdash T : K \quad \Gamma, id = T : K \vdash_{\delta} t' : T'}{\Gamma \vdash_{\delta} [id : K = T] \cdot t' : T'} \quad
\frac{\Gamma \vdash \{t' \simeq t'\} : \star}{\Gamma \vdash_{\Downarrow} \beta \{t\} : \{t' \simeq t'\}} \quad
\frac{\Gamma \vdash_{\delta} t : \{t_1 \simeq t_2\}}{\Gamma \vdash_{\delta} \varsigma t : \{t_2 \simeq t_1\}} \\
\\
\frac{\Gamma \vdash_{\Downarrow} t : \{|t_1| \simeq |t_2|\} \quad \Gamma \vdash_{\delta} t_1 : T}{\Gamma \vdash_{\delta} \phi t \cdot t_1 \{t_2\} : T} \quad
\frac{\Gamma \vdash_{\Downarrow} t : T}{\Gamma \vdash_{\Uparrow} \chi T \cdot t : T} \quad
\frac{\Gamma \vdash_{\Downarrow} t : \{\mathbf{tt} \simeq \mathbf{ff}\}}{\Gamma \vdash_{\Downarrow} \delta \cdot t : T} \quad 4
\end{array}$$

Figure 15: Type checking $\boxed{\Gamma \vdash_{\delta} s : C}$ (sans inductive datatypes)

The typing rule for ρ shows that ρ is a primitive for rewriting by an (untyped) equality. If t is an expression that synthesizes a proof that two terms t_1 and t_2 are equal, and t' is an expression synthesizing type $[t_1/x] T$ (where, as per the footnote, t_1 does not occur in T), then we may essentially rewrite its type to $[t_2/x] T$. The rule for β is reflexivity for equality – it witnesses that a term is equal to itself, provided that the type of the equality is well-formed. The rule for ς is symmetry for equality. Finally, ϕ acts as a “casting” primitive: the rule for its use says that if some term t witnesses that two terms t_1 and t_2 are equal, and t_1 has been judged to have type T , then intuitively t_2 can also be judged to have type T . (This intuition is justified by the erasure rule for ϕ – the expression erases to $|t_2|$). The last rule involving equality is for δ , which witnesses the logical principle *ex falso quodlibet* – if a certain impossible equation is proved (namely that the two Church-encoded booleans \mathbf{tt} and \mathbf{ff} are equal), then *any* type desired is inhabited. The remaining primitive χ allows the user to provide an explicit top-level annotation for a term.

5 Inductive Datatypes

Before we can provide the typing rules for introduction and usage of inductive datatypes, some auxiliary definitions must be given. The syntax for these, and the structure of this entire section, borrows heavily from

the conventions of the Coq documentation⁵. The author believes it is worthwhile to restate this development in terms of the Cedilleum type system, rather than merely pointing readers to the Coq documentation and asking them to infer the differences between the two systems.

To begin with, the production *defDataType* gives the concrete syntax for datatype definitions, but it is not a very useful notation for representing one in the abstract syntax tree. In our typing rules we will instead use the notation $\text{Ind}_M[p](\Gamma_I := \Sigma)$, where

- M is a meta-variable ranging over constant labels “C” and “A” (used to distinguish **con**crete and **ab**stracted inductive definitions – more on this below)
- p is the number of **p**arameters of the inductive definition
- Γ_I is a typing context binding *one* type variable I , the inductive type being defined
- Σ is a typing context containing the n data constructors c_1, \dots, c_n of I .

For example, consider the **List** and **Vec** definitions from Section 2. These will be represented in the AST as

$$\text{Ind}_C[1](\text{List} : \star \rightarrow \star := \begin{array}{ll} \text{nil} & : \forall A : \star. \text{List} \cdot A \\ \text{cons} & : \forall A : \star. A \rightarrow \text{List} \cdot A \rightarrow \text{List} \cdot A \end{array})$$

and

$$\text{Ind}_C[1](\text{Vec} : \star \rightarrow \text{Nat} \rightarrow \star := \begin{array}{ll} \text{vnil} & : \forall A : \star. \text{Vec} \cdot A \text{ zero} \\ \text{vcons} & : \forall A : \star. \forall n : \text{Nat}. A \rightarrow \text{Vec} \cdot A \ n \rightarrow \text{Vec} \cdot A \ (\text{succ } n) \end{array})$$

All inductive types the user will define will be concrete inductive definitions, and have global scope. Abstracted definitions are automatically generated during fix-point pattern matching, and have local scope.

For an inductive datatype definition to be well-formed, it must satisfy the following conditions (each of which is explained in more detail in Subsections 5.1 and 5.2):

- The kind of I must be (at least) a *p-arity of kind \star* .
- The types of each $id \in \Sigma$ must be *types of constructors of I*
- The definition must satisfy the *non-strict* positivity condition.

Similarly, the notation in the grammar of Cedilleum μ' and μ for pattern matching is inconvenient, and we will represent them in the AST as resp. $\mu'(t, P, t_{i=1..n})$ and $\mu(x_{\text{rec}}, I', x_{\text{to}}, t, P, t_{i=1..n})$. Translation from the form given in the grammar to this form is discussed in detail below, but is as expected. In particular, we enforce that patterns are exhaustive and non-overlapping, and that I' and x_{to} (which correspond to the automatically generated identifiers like **Nat/ih** and **fromNat/ih** from the introduction) are fresh w.r.t the global and local context. For example, consider the pattern-matches given in the code listings for **isvnil** and **vlength** above. These would be translated into the AST as

$$\mu'(xs, \Lambda n. \lambda x. \text{Bool}, \begin{array}{l} \text{tt} \\ \Lambda n. \lambda x. \lambda xs. \text{ff} \end{array})$$

and

$$\mu(\text{len}, \text{Vec}/\text{len}, \text{fromVec}/\text{len}, xs, \Lambda n. \lambda x. \text{Nat}, \begin{array}{l} \text{zero} \\ \Lambda n. \lambda x. \lambda xs. \text{succ } (\text{len} - n \text{ } xs) \end{array})$$

⁵<https://coq.inria.fr/refman/language/cic.html#inductive-definitions>

In general, the generated name for I' and x_{to} that users will write in Cedilleum programs will be of the form “ I/x_{rec} ” and “**from** I/x_{rec} ”.

For a pattern construct (μ or μ') in the AST to be well-formed, it must satisfy the following conditions (each of which is, again, explained in more detail in Subsections 5.3, 5.5, and 5.6):

- The motive P must be well-kinded
- P must be a legal motive to be used in eliminating the inductive type I of the scrutinee t
- Each branch t_i must have the type expected given the constructor $c_i \in \Sigma$ and the motive P .

5.1 Auxiliary Definitions

Contexts To ease the notational burden, we will introduce some conventions for writing contexts within terms and types.

- We write $\lambda \Gamma$, $\Lambda \Gamma$, $\forall \Gamma$, and $\Pi \Gamma$ to indicate some form of abstraction over each variable in Γ . For example, if $\Gamma = x_1 : T_1, x_2 : T_2$ then $\lambda \Gamma. t = \lambda x_1 : T_1. \lambda x_2 : T_2. t$. Additionally, we will also write $\prod \Gamma$ to indicate an arbitrary mixture of Π and \forall quantified variables. Note that *if $\prod \Gamma$ occurs multiple times within a definition or inference rule*, the intended interpretation is that *all occurrences have the same mixture of Π and \forall quantifiers*.
- $\|\Gamma\|$ denotes the length of Γ (the number of variables it binds)
- We write $s \Gamma$ to indicate the sequence of variable arguments in Γ given as arguments to s . Implicit in this notation is the removal of typing annotations from the variables Γ when these variables are given as arguments to s .

Since in Cedilleum there are three flavors of applications (to a type, to an erased term, and to an unerased term), we will only use this notion when the type or kind of s is known, which is sufficient to disambiguate the flavor of application intended for each particular binder in Γ . For example, if s has type $\forall X : \star. \forall x : X. \Pi x' : X. X$ and $\Gamma = X : \star, x : X, x' : X$ then $s \Gamma = s \cdot X \cdot x \cdot x'$

- Δ and Δ' are notations we will use for a specially designated contexts associating type variables with both global “concrete” and local “abstracted” inductive data-type declarations. The purpose of this latter sort of declaration is to enable type-guided termination of definitions using fixpoints (see Section 5.7) For example, given just the (global) data type declaration of Vec , we would have $\Delta(Vec) = \text{Ind}_C[1](\Gamma_{Vec} := \Sigma)$, where $\Gamma_{Vec} = Vec : \star \rightarrow Nat \rightarrow \star$ and Σ binds data constructors $vnil$ and $vcons$ to the appropriate types.

p -arity A kind K is a p -arity if it can be written as $\Pi \Gamma. K'$ for some Γ and K' , where $\|\Gamma\| = p$. For an inductive definition $\text{Ind}_M[p](\Gamma_I := \Sigma)$, requiring that the kind $\Gamma_I(I)$ is a p -arity of \star ensures that I *really does have* p parameters.

Types of Constructors T is a *type of a constructor of* I iff

- it is $I \ s_1 \dots s_n$
- it can be written as $\forall s : C. T$ or $\Pi s : C. T$, where (in either case) T is a type of a constructor of I

Positivity condition The positivity condition is defined in two parts: the positivity condition of a type T of a constructor of I , and the positive occurrence of I in T . We say that a type T of a constructor of I satisfies the positivity condition when

- T is $I\ s_1 \dots s_n$ and I does not occur anywhere in $s_1 \dots s_n$
- T is $\forall s:C. T'$ or $\Pi s:C. T'$, T' satisfies the positivity condition for I , and I occurs *only* positively in C

We say that I occurs only positively in T when

- I does not occur in T
- T is of the form $I\ s_1 \dots s_n$ and I does not occur in $s_1 \dots s_n$
- T is of the form $\forall s:C. T'$ or $\Pi s:C. T'$, I occurs only positively in T' , and I *does not* occur positively in C

5.2 Well-formed inductive definitions

Let Γ_P, Γ_I , and Σ be contexts such that Γ_I associates a single type-variable I to kind $\Pi \Gamma_P. K$ and Σ associates term variables $c_1 \dots c_n$ with corresponding types $\forall \Gamma_P. T_1, \dots \forall \Gamma_P. T_n$. Then the rule given in Figure 16 states when an inductive datatype definition may be introduced, provided that the following side conditions hold:

Figure 16: Introduction of inductive datatype

$$\frac{\emptyset \vdash \Gamma_I(I) : \square \quad \|\Gamma_P\| = p \quad (\Gamma_I, \Gamma_P \vdash T_i : \star)_{i=1..n}}{\text{Ind}_M[p](\Gamma_I := \Sigma) \text{ wf}}$$

- Names I and $c_1 \dots c_n$ are distinct from any other inductive datatype type or constructor names, and distinct amongst themselves
- Each of $T_1 \dots T_n$ is a type of constructor of I which satisfies the positivity condition for I . Furthermore, each occurrence of I in T_i is one which is applied to the parameters Γ_P .
- Identifiers I, c_1, \dots, c_n are fresh w.r.t the global context, and do not overlap with each other nor any identifiers in Γ_P .

When an inductive data-type has been defined using the *defDataType* production, it is understood that this always a concrete inductive type, and it (implicitly) adds to a global typing context the variable bindings in Γ_I and Σ . Similarly, when checking that the kind $\Gamma_I(I)$ and type T_i are well-sorted and well-kinded, we assume an (implicit) global context of previous definitions.

5.3 Valid Elimination Kind

Figure 17: Valid elimination kinds

$$\frac{}{\overline{\llbracket T : \star \mid T \rightarrow \star \rrbracket}} \quad \frac{\overline{\llbracket T\ s : K \mid K' \rrbracket}}{\overline{\llbracket T : \Pi s:C. K \mid \Pi s:C. K' \rrbracket}}$$

When type-checking a pattern match (either μ or μ'), we need to know that the given motive P has a kind K for which elimination of a term with some inductive data-type I is permissible. We write this judgment as $\llbracket T : K' \mid K \rrbracket$, which should be read “the type T of kind K' can be eliminated through pattern-matching with a motive of kind K ”. This judgment is defined by the simple rules in Figure 17. For example, a valid elimination kind for the indexed type family $Vec \cdot X$ (which has kind $\Pi n:Nat. \star$) is $\Pi n:Nat. \Pi x:Vec \cdot X\ n. \star$

5.4 Valid Branch Type

Another piece of kit we need is a way to ensure that, in a pattern-matching expression, a particular branch has the correct type given a particular constructor of an inductive data-type and a motive. We write $\{\{c : T\}\}_I^P$ to indicate the type corresponding to the (possibly partially applied) constructor c of I and its type T . We abbreviate this notation to $\{\{c\}\}^P$ when the inductive type variable I , and the type T of c , is known from the (meta-language) context.

$$\begin{aligned} \{\{c : I \bar{T} \bar{s}\}\}_I^P &= P \bar{s} c \\ \{\{c : \forall x : T'. T\}\}_I^P &= \forall x : T'. \{\{c \cdot x : T\}\}_I^P \\ \{\{c : \forall x : K. T\}\}_I^P &= \forall x : K. \{\{c \cdot x : T\}\}_I^P \\ \{\{c : \Pi x : T'. T\}\}_I^P &= \Pi x : T'. \{\{c x : T\}\}_I^P \end{aligned}$$

where we leave implicit the book-keeping required to separate the parameters \bar{T} from the indices \bar{s} .

The biggest difference between this definition and the similar one found in the Coq documentation is that types can have implicit and explicit quantifiers, so we must make sure that the types of branches have implicit / explicit quantifiers (and the subjects c have applications for types, implicit terms, and explicit terms), corresponding to those of the arguments to the data constructor for the pattern for the branch.

5.5 Well-formed Patterns

Figure 18: Well-formedness of a pattern

$$\frac{\Gamma \vdash P : K \quad \Sigma = c_1 : \forall \Gamma_P. T_1, \dots, c_n : \forall \Gamma_P. T_n \quad \|\bar{T}\| = \|\Gamma_P\| = p \quad \llbracket I \bar{T} : \Gamma(I) \mid K \rrbracket \quad (\Gamma, \Delta \vdash_{\downarrow} t_i : \{\{c_i \bar{T}\}\}^P)_{i=1..n}}{WF\text{-}Pat(\Gamma, \Delta, \text{Ind}_M[p](\Gamma_I := \Sigma), \bar{T}, \mu'(t, P, t_{i=1..n}))}$$

Figure 18 gives the rule for checking that a pattern $\mu'(t, P, t_{i=1..n})$ is well-formed. We check that the motive P is well-kinded at kind K , that the given parameters \bar{T} match the expected number p from the inductive data-type declaration, that an inductive data-type I instantiated with the given parameters \bar{T} can be eliminated to a type of kind K , and that the given branches t_i account for each of the constructors c_i of Σ and have the required branch type $\{\{c_i \bar{T}\}\}^P$ under the given local context Γ and context of inductive data-type declarations Δ .

5.6 Generation of Abstracted Inductive Definitions

Cedilleum supports *histomorphic* recursion (that is, having access to all previous recursive values) where termination is ensured through typing. In order to make this possible, we need a mechanism for tracking the global definitions of *concrete* inductive data types as well the locally-introduced *abstract* inductive data type representing the recursive occurrences suitable for a fixpoint function to be called on.

If I is an inductive type such that $\Delta(I) = \text{Ind}_C[p](\Gamma_I := \Sigma)$ and I' is a fresh type variable, then we define function $Hist(\Delta, I, \bar{T}, I')$ producing an abstracted (well-formed) inductive definition $\text{Ind}_A[0](\Gamma_{I'} := \Sigma')$, where

- $\Gamma_{I'}(I') = \forall \Gamma_D. \star$ if $\Gamma_I(I) = \forall \Gamma_P. \forall \Gamma_D. \star$ (and $\|\Gamma_P\| = \|\bar{T}\| = p$)

That is, the kind of I' is the same as the kind of $I \bar{T}$

- $\Sigma' = c'_1 : \forall \Gamma_D. \prod \Gamma_{A'_1}. I' \Gamma_D, \dots, c'_n : \forall \Gamma_D. \prod \Gamma_{A'_n}. I \bar{T} \Gamma_D,$

when each of the concrete constructors c_i in Σ are associated with type $\forall \Gamma_P. \forall \Gamma_D. \prod \Gamma_{A_i}. I \Gamma_P \Gamma_D$ and each $\Gamma_{A'_i} = [\lambda \Gamma_P. I' / I, \bar{T} / \Gamma_P] \Gamma_{A_i}$.

That is, transforming the concrete constructors of the inductive datatype I to “abstracted” constructors involves replacing each recursive occurrence of $I \Gamma_P$ with the fresh type variable I , and instantiating each of the parameters Γ_P with \bar{T} .

Users of Cedilleum will see “punning” of the concrete constructors c_i and abstracted constructors c'_i . In particular, when using fix-point pattern matching branch labels will be written with the constructors for the concrete inductive data-type, and the expected type of a branch given by the motive will pretty-print using the concrete constructors. In the inference rules, however, we will take more care to distinguish the abstract constructors (see Subsection 5.7).

5.7 Typing Rules

Figure 19: Use of an inductive datatype $\text{Ind}_M[p](\Gamma_I := \Sigma)$

$$\begin{array}{c}
\frac{\Gamma \vdash_{\uparrow} t : I \bar{T} \bar{s} \quad \text{WFPat}(\Gamma, \Delta, \Delta(I), \bar{T}, \mu'(t, P, t_{i=1..n}))}{\Gamma, \Delta \vdash_{\delta} \mu'(t, P, t_{i=1..n}) : P \bar{s} t} \\
\\
\begin{array}{l}
\Gamma \vdash_{\uparrow} t : I \bar{T} \bar{s} \quad \Delta(I) = \text{Ind}_C[p](\Gamma_I := \Sigma) \quad \Gamma_I(I) = \Pi \Gamma_P. \Pi \Gamma_D. \star, \|\Gamma_P\| = p \quad \text{Hist}(\Delta, I, \bar{T}, I') = \text{Ind}_A[0](\Gamma_{I'} := \Sigma') \\
\Gamma' = \Gamma, \Gamma_{I'}, x_{\text{to}} = \Lambda \Gamma_D. \lambda x. x : \forall \Gamma_D. I' \Gamma_D \rightarrow I \bar{T} \Gamma_D, x_{\text{rec}} : \forall \Gamma_D. \Pi x : I' \Gamma_D. P \Gamma_D (x_{\text{to}} \Gamma_D x) \quad \Delta' = \Delta, \text{Hist}(\Delta, I, \bar{T}, I') \\
\text{WFPat}(\Gamma', \Delta', \Delta'(I'), \emptyset, \mu'(t, P, t_{i=1..n})) \\
\hline
\Gamma, \Delta \vdash_{\delta} \mu(x_{\text{rec}}, I', x_{\text{to}}, t, P, t_{i=1..n}) : P \bar{s} t
\end{array}
\end{array}$$

The first rule of Figure 19 is for typing simple pattern matching with μ' . We need to know that the scrutinee t is well-typed at some inductive type $I \bar{T} \bar{s}$, where \bar{T} represents the parameters and \bar{s} the indices. Then we defer to the judgment WFPat to ensure that this pattern-matching expression is a valid elimination of t to type P .

The second rule is for typing pattern-matching with fix-points, and is significantly more involved. As above we check the scrutinee t has some inductive type $I \bar{T} \bar{s}$. We confirm that I is a *concrete* inductive data-type by looking up its definition in Δ , and then generate the abstracted definition $\text{Hist}(\Delta, I, \bar{T}, I')$ for some fresh I' . We then add to the local typing context $\Gamma_{I'}$ (the new inductive type I' with its associated kind) and two new variables x_{to} and x_{rec} .

- x_{to} is the *revealer*. It casts a term of an abstracted inductive data-type $I' \Gamma_D$ to the concrete type $I \bar{T} \Gamma_D$. Crucially, it is an *identity* cast (the implicit quantification $\Lambda \Gamma_D$ disappears after erasure). The intuition why this should be the case is that the abstracted type I' only serves to mark the recursive occurrences of I during pattern-matching to guarantee termination.
- x_{rec} is the *recursor* (or the inductive hypothesis). Its result type $P' \Gamma_D x$ utilizes x_{to} in P' to be well-typed, as the x in this expression has type $I' \Gamma_D$, but P expects an $I \bar{T} \Gamma_D$. Because x_{to} erases to the identity, uses of the x_{rec} will produce expressions whose types will not interfere with producing the needed result for a given branch (see the extended example – TODO).

With these definitions, we finish the rule by checking that the pattern is well-formed using the augmented local context Γ' and context of inductive data-type definitions Δ' .

6 Elaboration of Inductive Datatypes

As mentioned in Section 1, Cedilleum is not based on CIC. Rather, its core theory is the *Calculus of Dependent Lambda Eliminations* (CDLE), whose complete typing rules can be those of Section 4 plus rules for dependent

intersections (see [Stu18]). That is to say, the preceding treatment for inductive datatypes (Section 5) is a high-level and convenient interface for *derivable* inductive λ -encodings. This section explains the elaboration process. Since the generic derivation of inductive data-types with course-of-value induction has been covered in-depth in [TODO], we omit these details and instead describe the *interface* such developments provide which data-type elaboration targets.

At a high level, inductive data-types in Cedilleum are first translated to *identity mappings*, which are (in the non-indexed case) a class of type schemes $F: \star \rightarrow \star$ that are more general than functors. The parameter of the identity scheme replaces all recursive occurrences of the data-type in the signatures of the constructor and a quantified type variable replaces all “return type” occurrences. For example, the type scheme for data-type `Nat` is $\lambda R: \star. \forall X: \star. X \rightarrow (R \rightarrow X) \rightarrow X$, with R the parameter and X the quantified variable. For the rest of this section we assume the reader has at least a basic understanding of impredicative encodings of datatypes (see [PPM89] and [Wad90]) and taking the least fix-point of functors (see [MFP91]).

The following developments are parameterized by an indexed type scheme F of kind $(\Pi \Gamma_D. \star) \rightarrow (\Pi \Gamma_D. \star)$ corresponding to the kind $\Pi \Gamma_D. \star$ of inductive data-type I declared as $\text{Ind}_I[p](\Gamma_I := \Sigma)$

6.1 Identity Mappings

Our first task is to describe identity mappings, the class of type schemes $F: (\Pi \Gamma_D. \star) \rightarrow \Pi \Gamma_D. \star$ we concerned with. Identity mappings are similar to functors in that they come equipped with a function that resembles `fmap`: $\forall \Gamma_D. \forall A B: \Pi \Gamma_D. \star. \Pi f: (A \cdot \Gamma_D \rightarrow B \cdot \Gamma_D). F \cdot (A \cdot \Gamma_D) \rightarrow F \cdot (B \cdot \Gamma_D)$ except that it need only be defined for an argument f that is equal to the identity function. We define the type `Id` of such functions and declare (indicated by $\langle \dots \rangle$) its elimination principle `elimIdD`:

$\text{Id}_D : \Pi A B: (\Pi \Gamma_D. \star). \iota \text{id}: \forall \Gamma_D. A \Gamma_D \rightarrow B \Gamma_D. \{\text{id} \simeq \lambda x. x\}.$
 $\text{elimId}_D : \forall A B: (\Gamma_D. \star). \text{Id}_D \cdot A \cdot B \Rightarrow A \rightarrow B = \langle \dots \rangle$

Recall that since Cedilleum has a Curry-style type system and implicit products there are many non-trivial functions that erase to identity. While the definition of `elimIdD` is omitted, it is important to note that it enjoys the property of erasing to the identity function:

$\text{elimId}_D\text{-prop} : \{\text{elimId}_D \simeq \lambda x. x\} = \beta.$

We may now define `IdMapping` as a scheme F that comes with a way to lift identity functions:

$\text{IdMapping}_D : \Pi F: (\Gamma_D \rightarrow \star) \rightarrow (\Gamma_D \rightarrow \star). \star$
 $= \lambda F. \forall A B: (\Gamma_D \rightarrow \star). \Pi_{\Gamma_D} \text{Id}_D \cdot A \cdot B \rightarrow \text{Id}_D \cdot (F \cdot A) \cdot (F \cdot B).$

Finally, it is convenient to define `fimap` which given an `IdMapping` and an `Id` function performs the lifting:

$\text{fimap}_D : \forall F: (\Pi \Gamma_D. \star) \rightarrow (\Pi \Gamma_D. \star). \forall \text{id}: \text{IdMapping}_D \cdot F. \text{Cast}_D \cdot A \cdot B \Rightarrow F \cdot A \rightarrow F \cdot B$
 $= \Lambda F \text{id} c. \lambda f. \text{elimId}_D \text{-(id c) f}.$

From `elimIdD-prop` it should be clear that `fimapD` also erases to $\lambda x. x$.

6.2 Type-views of Terms

A crucial component of course-of-value is the ability to view some term as having two different types. The idea behind a `View` is similar to that behind the type `Id` from the previous section, except now we explicitly name the doubly-typed term:

$\text{View} : \Pi A: \star. A \rightarrow \star \rightarrow \star = \lambda A a B. \iota b: B. \{a \simeq b\}$
 $\text{elimView} : \forall A B: \star. \Pi a: A. \text{View} \cdot A \cdot a \cdot B \Rightarrow B = \langle \dots \rangle$
 $\text{elimView-prop} : \{\text{elimView} \simeq \lambda x. x\} = \beta.$

6.3 λ -encoding Interface

This subsection describes the interface to which data-type declarations are elaborated; it is parameterized by an identity mapping.

`module (FD: (Π ΓD. ★) → (Π ΓD. ★)){im: IdMapping ·FD}.`

where parameters F_D and im are automatically derived from the declaration of a positive data-type.

With these two parameters alone, the generic developments of [TODO] provide the following interface for inductive λ -encodings of data-types:

```

FixD : Π ΓD. ★ = <..>
inD  : ∀ ΓD. FD ·FixD ΓD → FixD ΓD = <..>
outD : ∀ ΓD. FixD ΓD → FD ·FixD ΓD = <..>

PrfAlgD : Π P: (Π ΓD. Π d: FixD ΓD. ★). ★
= λ P. ∀ R: (Π ΓD. ★).
  ∀ c: IdD ·R ·FixD.
  Π v: View ·(∀ ΓD. FixD ΓD → FD ·FixD ΓD) out ·(∀ ΓD. R ΓD → FD ·R ΓD).
  Π ih: (∀ ΓD. Π r: R ΓD. P ΓD (elimIdD -c -ΓD r)).
  Π ΓD. Π fr: F ·R ΓD.
  P ΓD (inD -ΓD (fimapD -im -c fr)).
inductionD : ∀ P: (Π ΓD. Π d: FixD ΓD. ★). PrfAlgD ·P → ∀ ΓD. Π d: FixD ΓD. P ΓD d
= <..>

```

The first three definitions give Fix_D as the (least) fixed-point of F_D, with in_D and out_D representing resp. a generic set of constructors and destructors. induction_D of course is the proof-principle stating that if one can provide a PrfAlg for property P (that is, P holds for all Fix_D generated by (generic) constructor in_D) then this suffices to show that P holds for *all* Fix_D.

We now explain the definition of PrfAlg_D in more detail:

- R is the type of recursive occurrences of the data-type Fix_D.

It corresponds directly to types like `rec/Nat` when using μ in Cedilleum

- c is a “revealer”, that is to say a proof that R really *is* Fix_D witnessed by an identity function.

It corresponds directly to functions like `rec/cast` when using μ

- v is evidence that the (generic) destructor out_D can be used on the recursive occurrence type R for further pattern-matching.

It corresponds directly to μ' (when used outside of μ it corresponds to the “trivial” view that out_D has the type it is already declared to have).

- ih is the inductive hypothesis, stating that property P holds for all recursive occurrences R of an inductive case

It corresponds directly to the μ -bound variable for fix-point recursion.

- fr represents the collection of constructors that each μ branch must account for.

For example, for the data-type Nat we have identity mapping `fr: ∀ X: ★. X → (R → X) → X` and Cedilleum cases branches `{ | zero → zcase | succ r → scase r }` translate to `fr zcase (λ r. scase r)`

- Finally, result type P Γ_D (in_D -Γ_D (fimap_D -im -c fr)) accounts for the return type of each case branch.

Since P is phrased over Fix_D, and we have by assumption `fr: FD ·R ΓD`, we must first use our identity mapping im to traverse fr and cast each recursive occurrence R Γ_D to Fix_D Γ_D, producing an expression of type F ·Fix_D Γ_D which we are then able to transform into Fix_D Γ_D using (generic) constructor in_D.

While the definitions of in_D , out_D , and induction_D are omitted, it is important that they have the following computational behavior (guaranteed by [TODO]):

```
lambek1D : ∀ ΓD. Π gr: FD FixD ΓD. {outD (inD gr) ≈ gr} = β.
lambek2D : ∀ ΓD. Π d: FixD ΓD. {in (out d) ≈ d}
  = inductionD · (λ ΓD. λ x: FixD ΓD. {in (out x) ≈ x})
    (Λ R. Λ c. λ o. Λ eq. λ ih. λ gr. β).
```

```
inductionCancelD : ∀ P: (Π ΓD. FixD ΓD → ★).
  Π alg: PrfAlg · P → ∀ ΓD. Π fr: F · FixD ΓD.
  { inductionD alg (in gr) ≈ alg outD (inductionD alg) fr }
  = λ _ . λ _ . β.
```

That is, in_D and out_D are inverses of each other and induction_D behaves like a fold (where the algebra takes the additional out_D argument).

6.4 Sum-of-Products Induction

As stated above, every inductive data-type declaration $\text{Ind}_I[p](\Gamma_I := \Sigma)$ is first translated to a type-scheme IF where all recursive occurrences of type I in the constructor signatures Σ have been replaced by the scheme's argument R. In this subsection describe that process more precisely and explain “sum-of-products” induction for IF

First, as the kind of I is $\Pi \Gamma_p. \Pi \Gamma_D. \star$, where Γ_p are the parameters and Γ_D the indices, it follows that the kind of IF is $\Pi \Gamma_p. \Pi R: (\Pi \Gamma_D. \star). (\Pi \Gamma_D. \star)$. Next, each constructor c_j has type $\Sigma(c_j)$ which we know has the form $\prod \Gamma_j. I \Gamma_p \overline{t_j}$ (that is, some number of arguments Γ_j with a return type constructing the inductive data-type I). All recursive occurrences of I in Γ_j are substituted away with $\lambda \Gamma_p. R$ to produce Γ_j^R . With that, we may defined IF as

$$\lambda \Gamma_p R \Gamma_D. \forall X : \Pi \Gamma_D. \star. (\Pi c_j : (\prod \Gamma_j^R. X \overline{t_j}))_{j=1..n}. X \Gamma_D$$

Example The data-type declaration of **Vec** translates to:

```
VecF : Π A: ★. (Nat → ★) → Nat → ★
  = λ A R n. ∀ X: Nat → ★. X zero → (∀ n: Nat. A → R n → X (succ n)) → X n.
```

An induction principle for each of these non-recursive sum-of-products types IF can be defined in an automated way following the recipe given by [TODO]; in general these have the following shape:

```
indIF : ∀ Γp. ∀ R: (Π ΓD. ★). ∀ ΓD. Π fr: IF Γp · R ΓD. ∀ P: (Π ΓD. IF Γp · R ΓD → ★)
  (Π pj: ∏ ΓjR. P (cj ΓjR))j=1..n. P ΓD fr = <...>
```

A Deriving IdMapping_D for a Data-type Type Scheme

A type scheme F derived from a data-type declaration has by assumption a definition following the pattern:

```
F : Π Γp. (Π ΓD. ★) → Π ΓD. ★
  = λ Γp R ΓD. ∀ X: (Π ΓD. ★). (Π cj: (∏ ΓjR. X  $\overline{t_j}$ ))j=1..n. X ΓD
```

where R occurs only positively. From this we must give a witness that F is an identity mapping over R

```
idmap : ∀ Γp. IdMappingD · F
  = Λ Γp. Λ R1. Λ R2. Λ id. •
```

where the expected type of \bullet is $\text{Id}_D \cdot (F \cdot \Gamma_P \text{ R1}) \cdot (F \cdot \Gamma \text{ R2})$

We refine \bullet by the introduction rule for intersections (which Id_D is) and introduce the assumption $\text{fr1}: F \cdot \Gamma_P \text{ R1} \cdot \Gamma_D$

$[\Lambda \Gamma_D. \lambda \text{ fr1}. \bullet_1, \bullet_2]$

where $\bullet_1: F \cdot \Gamma_P \text{ R2} \cdot \Gamma_D$ and $\bullet_2: \{\lambda \text{ fr1}. \bullet_1 \simeq \lambda x. x\}$. As the only (non-hole) refinements we will make to \bullet_1 are converting terms to η -long form and applying $\text{elimId}_D \text{ -id}$ to subterms (which reduces to the identity function), we are justified in replacing \bullet_2 with β . We now refine the remaining \bullet_1 to

$\Lambda X. \lambda \bar{c}. \bullet \text{ fr1 } \bar{c}$

where each abstract constructor c_j in \bar{c} has type $\prod_{\forall} \Gamma^{\text{R2}}_j \cdot X \bar{c}_j$. Note again the superscript R2 – we are now trying to construct a term of type $F \cdot \Gamma_P \text{ R2} \cdot \Gamma_D$ so we assume the “abstract” constructors whose recursive occurrence types are R2 . Correspondingly, this means that $\bullet: F \cdot \Gamma_P \text{ R1} \cdot \Gamma_D \rightarrow (\prod c_j: (\prod_{\forall} \Gamma^{\text{R2}}_j \cdot X \bar{c}_j))_{j=1-n} \rightarrow X \Gamma_D$.

Since fr1 produces a value of type $X \Gamma_D$ when fed appropriate arguments, we refine \bullet by n holes \bullet_j applied to constructor c_j . The expression $\bullet \text{ fr1 } \bar{c}$ becomes

$\text{fr1 } (\bullet_j c_j)_{j=1-n}$

where now $\bullet_j: (\prod_{\forall} \Gamma^{\text{R2}}_j \cdot X \bar{c}_j) \rightarrow \prod_{\forall} \Gamma^{\text{R1}}_j \cdot X \bar{c}_j$. We henceforth dispense with the subscript j numbering the constructor and treat each abstract constructor uniformly.

A.1 Conversion of the Abstract constructors

We first make the expression $\bullet \text{ c}$ η -long, as in $\lambda_{\Lambda} \Gamma^{\text{R1}}. \bullet \text{ c } \Gamma^{\text{R1}}$, then refine $\bullet \text{ c } \Gamma^{\text{R1}}$ to an expression with m holes \bullet_k for each $y_k \in \Gamma^{\text{R1}}$ (where $m = \|\Gamma^{\text{R1}}\|$), yielding

$\text{c } (\bullet_k y_k)_{k=1-m}$

where $\bullet_k: \Gamma^{\text{R1}}(y_k) \rightarrow \Gamma^{\text{R2}}_k(y_k)$ (and the type of y_k and $\bullet_k y_k$ can depend resp. on any y^{R1}_j and $\bullet_j y_j$ where $j < k$). We now dispense with the subscript k for arguments and handle each constructor sub-data uniformly.

A.2 Conversion of Constructor Sub-data With Positive Recursive Occurrences

We now consider $\bullet \text{ y}$ where $y: S$ is some sub-data to an (abstract) constructor with recursive occurrence type R1 passing the positivity checker. (The expression $\bullet \text{ y}$ has type $[\text{R2}/\text{R1}]S$). There are two cases to consider:

- 1 R1 does not occur in the type of y

Refine \bullet to $\text{unit}: \forall X. \star. X \rightarrow X = \Lambda X. \lambda x. x$ and finish.

- 2 R1 occurs positively in the type of y

This means S has the shape $\prod_{\forall} \Gamma^{\text{R1}}_x \cdot T$ (where T is not formed by an arrow) with R1 occurring *only negatively* in the type of the $x_j \in \Gamma^{\text{R1}}_x$ (where $j = 1.. \|\Gamma^{\text{R1}}_x\|$). Make $\bullet \text{ y}$ η -long and refine the expression to $\|\Gamma^{\text{R1}}_x\|$ holes \bullet_j such that the expression is now

$\lambda_{\Lambda} \Gamma^{\text{R2}}_x. \bullet \text{ y } (\bullet_j x_j)_{j=1-n}$

Where here x_j is bound by Γ^{R2} and thus has negative occurrences of R2 . Note that we still require \bullet since it might be the case that $T = \text{R1 } \Gamma_D$ (handled below); it has type $S \rightarrow \prod_{\forall} \Gamma^{\text{R1}}_x \cdot [\text{R1}/\text{R2}]T$. Each \bullet_j has type $\Gamma^{\text{R2}}_x(x_j) \rightarrow \Gamma^{\text{R1}}_x(x_j)$.

Perform the steps outlined in Section A.3 to fill in each \bullet_j producing from $\bullet_j x_j$ the sequence of arguments \bar{c}_j of type Γ^{R1}_x that erase to $x_{j=1-n}$. Finally, refine \bullet to either unit or $\lambda y. \lambda x_j. \text{elimId } \text{-c } (y \text{ } x_j)$ depending on whether $T = \text{R1 } \Gamma_D$

A.3 Conversion of Constructor Sub-data With Negative Recursive Occurences

We consider $\bullet \ x$ where $x : \prod_y \Gamma_y^{R2}$. S , S is not an arrow and does not contain $R2$, and $R2$ occurs positively in the types of the variables bound by Γ_y^{R2} . The expression $\bullet \ x$ has type $\prod_y \Gamma_y^{R1}$. S .

Make $\bullet \ x$ η -long and introduce holes \bullet_j to apply to the sub-data as in

$$\lambda_{\Lambda} \Gamma_y^{R1} \bullet \ x \ (\bullet_j \ y_j)_{j=1-n}$$

where $\bullet_j : \Gamma_y^{R1}(y_j) \rightarrow \Gamma_y^{R2}(y_j)$. Perform the steps outlined by Section A.2 to fill in each \bullet_j producing from $\bullet_j \ y_j$ the sequence of arguments \bar{v} that erase to $y_{j=1-n}$.

References

- [Inr18] Inria. The Coq Documentation. <https://coq.inria.fr/refman/index.html>, 2018.
- [MFP91] Erik Meijer, Maarten Fokkinga, and Ross Paterson. Functional programming with bananas, lenses, envelopes and barbed wire. In *Conference on Functional Programming Languages and Computer Architecture*, pages 124–144. Springer, 1991.
- [Miq01] Alexandre Miquel. The implicit calculus of constructions: Extending pure type systems with an intersection type binder and subtyping. In *Proceedings of the 5th International Conference on Typed Lambda Calculi and Applications*, TLCA’01, pages 344–359, Berlin, Heidelberg, 2001. Springer-Verlag.
- [PM15] Christine Paulin-Mohring. Introduction to the calculus of inductive constructions, 2015.
- [PPM89] Frank Pfenning and Christine Paulin-Mohring. Inductively defined types in the calculus of constructions. In *International Conference on Mathematical Foundations of Programming Semantics*, pages 209–228. Springer, 1989.
- [Stu18] Aaron Stump. Syntax and semantics of cedille, 2018.
- [Wad90] Philip Wadler. Recursive types for free!, 1990.